

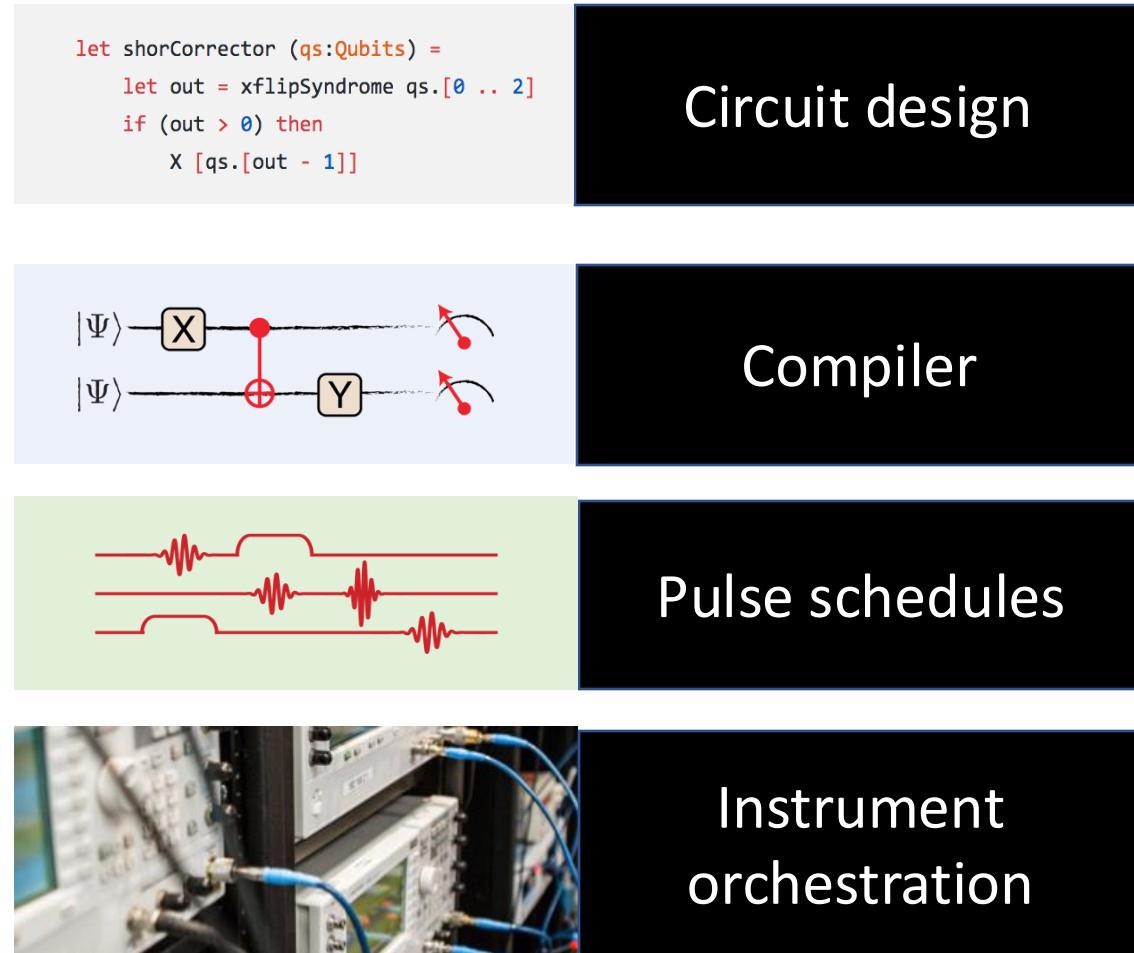
SW Stack for NISQ devices

Dr. Miroslav Dobsicek

Presentation overview

- ❖ SW stack overview
- ❖ User-space quantum stack
- ❖ Circuit level assembly
- ❖ Hardware level encoding

SW stack

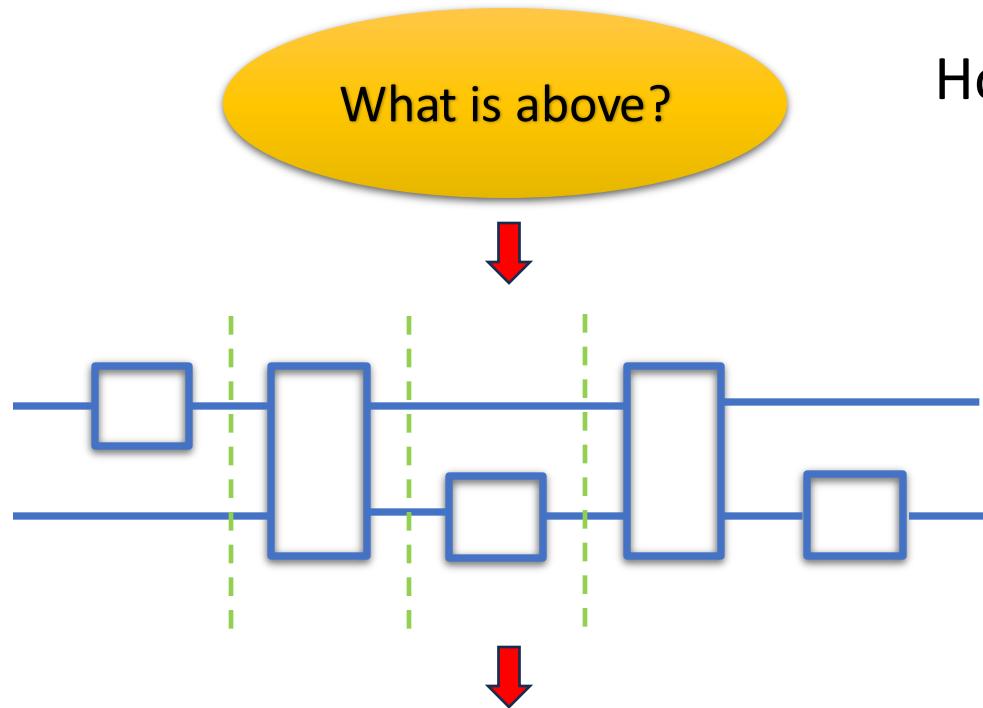


Computer science domain
Output for idealized quantum computer

Co-design for NISQ devices

Experimentalist domain
Single-user environment, lab work

SW stack is built around the quantum circuit model



How do we get a circuit?

What is above?

How do we run it?



High level parts of a SW stack

How do we generate quantum circuits?

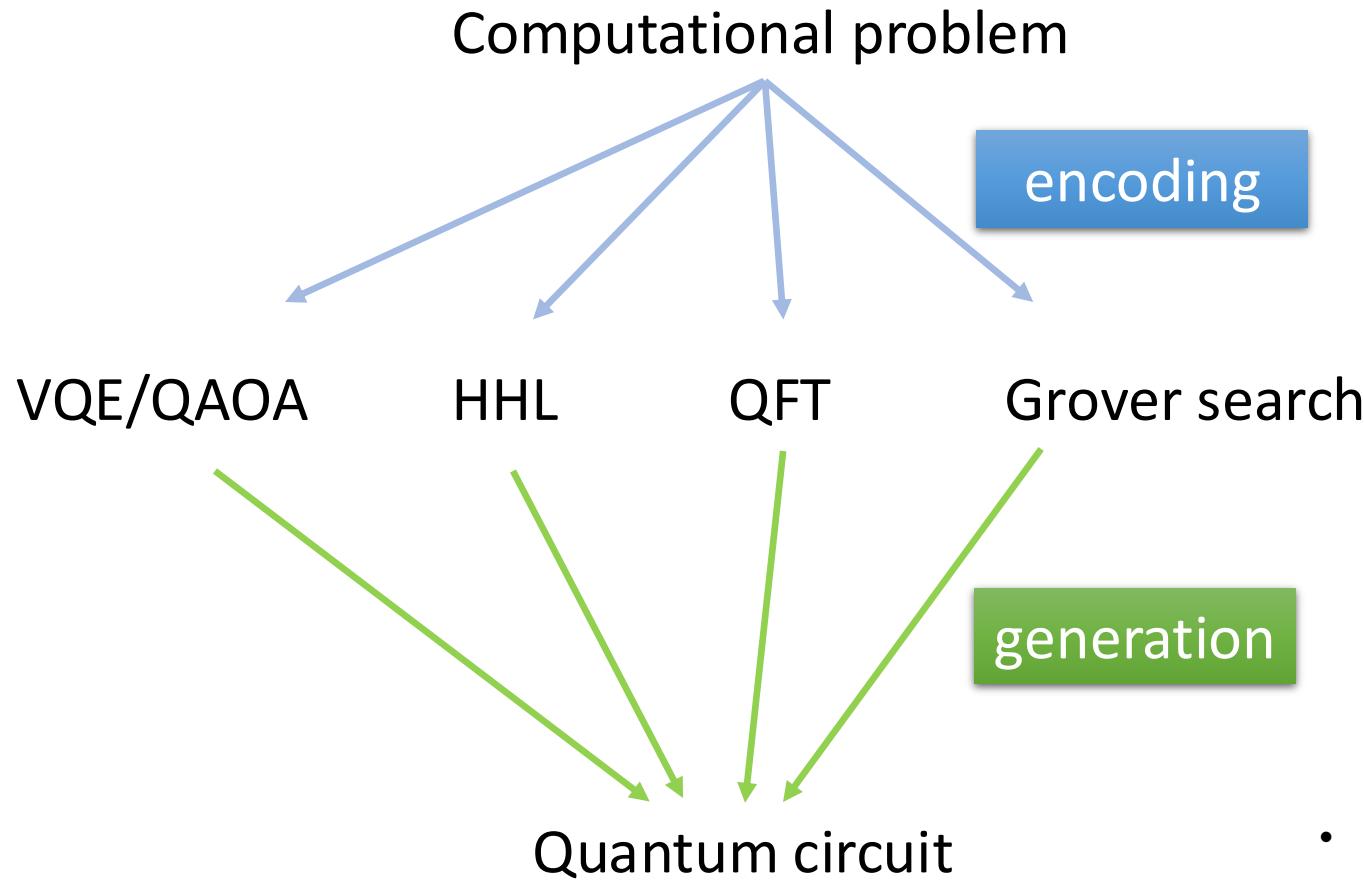
Generic methods

- ❖ **Encode** your problem into known quantum algorithms
- ❖ **Embed** a classical circuit into a quantum one through reversible logic
- ❖ **Automatically decompose** large transformations into sequences of smaller ones

Attacking directly the problem

- ❖ **Design** your own quantum algorithm

1. Problem encoding into an existing quantum algorithm

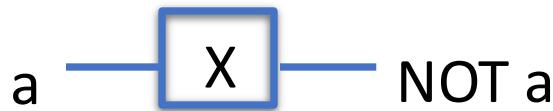
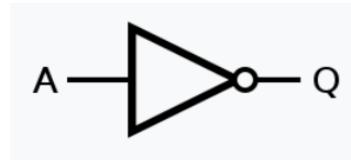


This is currently the most feasible way how to do a computation on a quantum computer.

- VQE – quantum chemistry problems
- QAOA – combinatorial opt. problems
- HHL – systems of linear equations (ML)
- QFT – detect group-like properties
- Grover search – generic square root speed-up

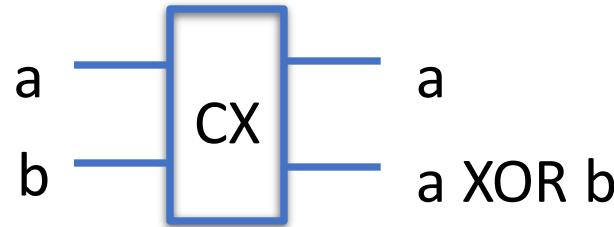
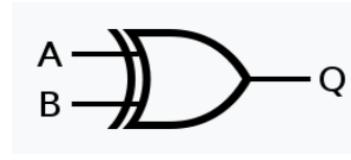
2. Embedding of classical circuits via reversible logic

NOT



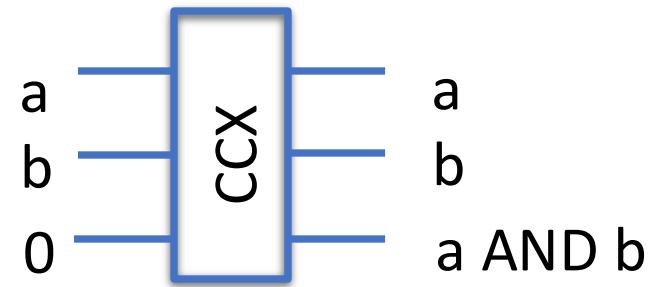
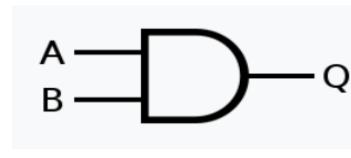
Classical logical gates mostly map to quantum gates in 1:1 fashion.

XOR



A quantum circuit generated in this way will have the same **overall** complexity as the classical circuit. Not better or worse. But, it will be capable of working with superpositions of states!

AND

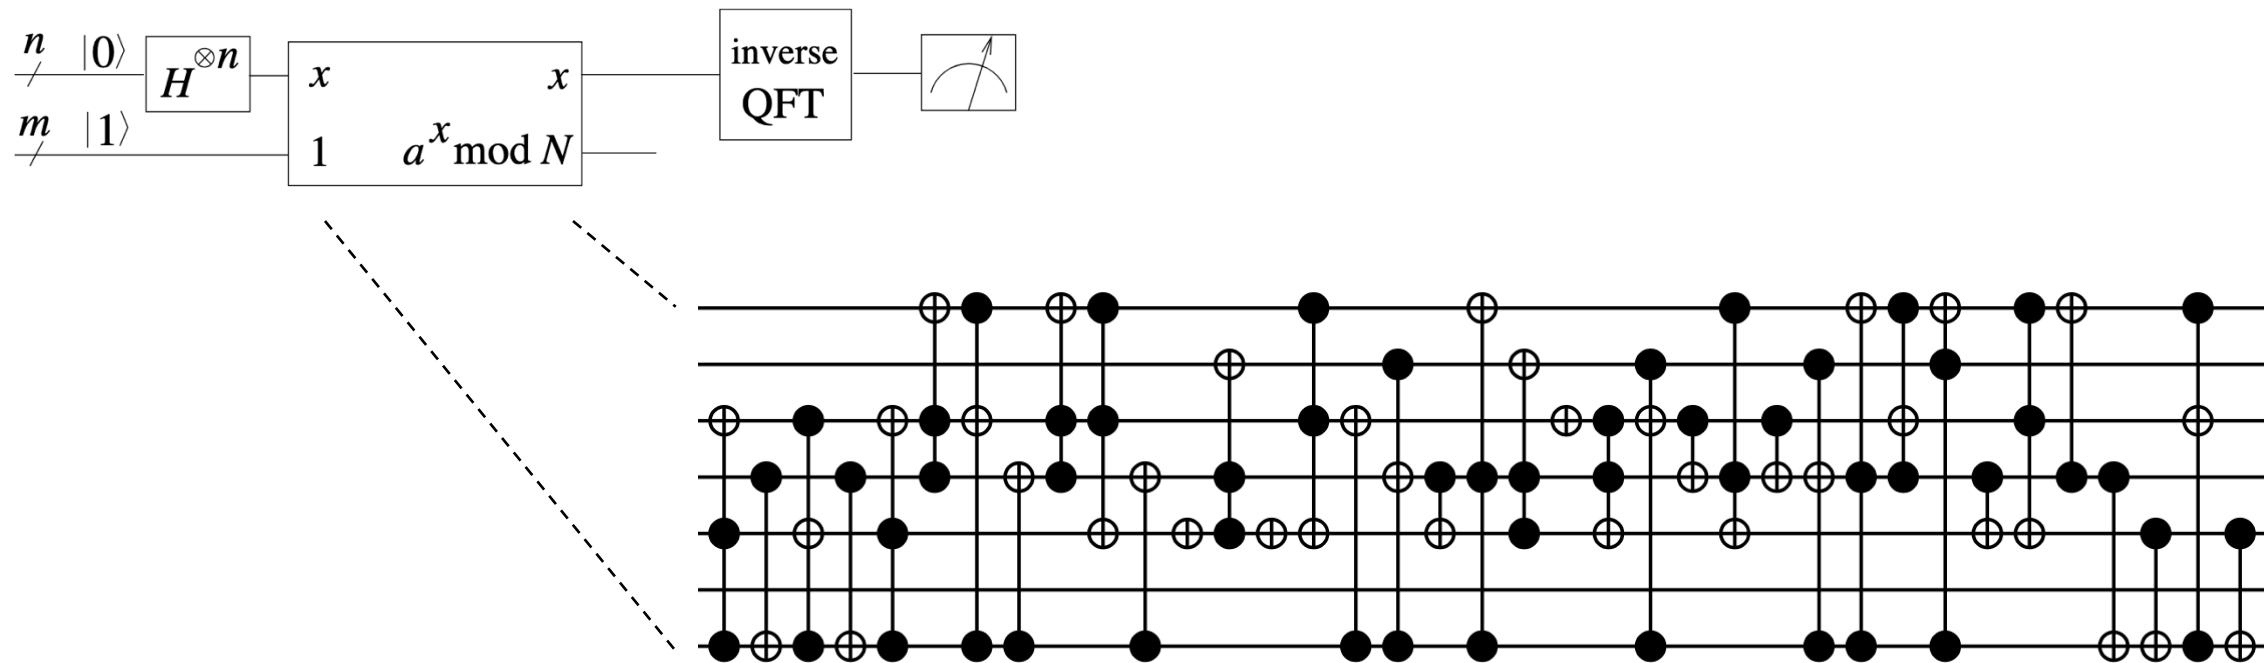


The cost are extra qubits guaranteeing reversibility.

Do you know that the QFT circuit and the circuit for a classical FFT are structurally the same?

2. Embedding of classical circuits via reversible logic

Part of the modular exponentiation circuit in Shor's algorithm generated by an embedding of a classical circuit.



Blog post: **Why haven't quantum computers factored 21 yet?**

<https://algassert.com/post/2500>

3. Automatic decomposition

Integers: factorization to prime numbers

$$12 = 2^2 \times 3$$

Matrices: decomposition to singular values

$$\begin{pmatrix} 2 & 1 \\ 1 & 2 \end{pmatrix} = \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & 1 \\ 1 & -1 \end{pmatrix} \begin{pmatrix} 3 & 0 \\ 0 & 1 \end{pmatrix} \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & 1 \\ 1 & -1 \end{pmatrix}$$

You start with a mathematical description of the desired unitary transformation and write it down in a matrix form. Then apply unitary decomposition algorithm(s). This process is usually based on **Singular Value Decomposition (SVD)**.

This approach is unlikely to lead to efficient circuits! The number of generated gates is **generally exponential** in the number of qubits.

Mathematical transformation

Such as $f : x \mapsto x^2$, or

$$\text{QFT} : |x\rangle \mapsto \frac{1}{\sqrt{N}} \sum_{k=0}^{N-1} \omega_N^{xk} |k\rangle.$$

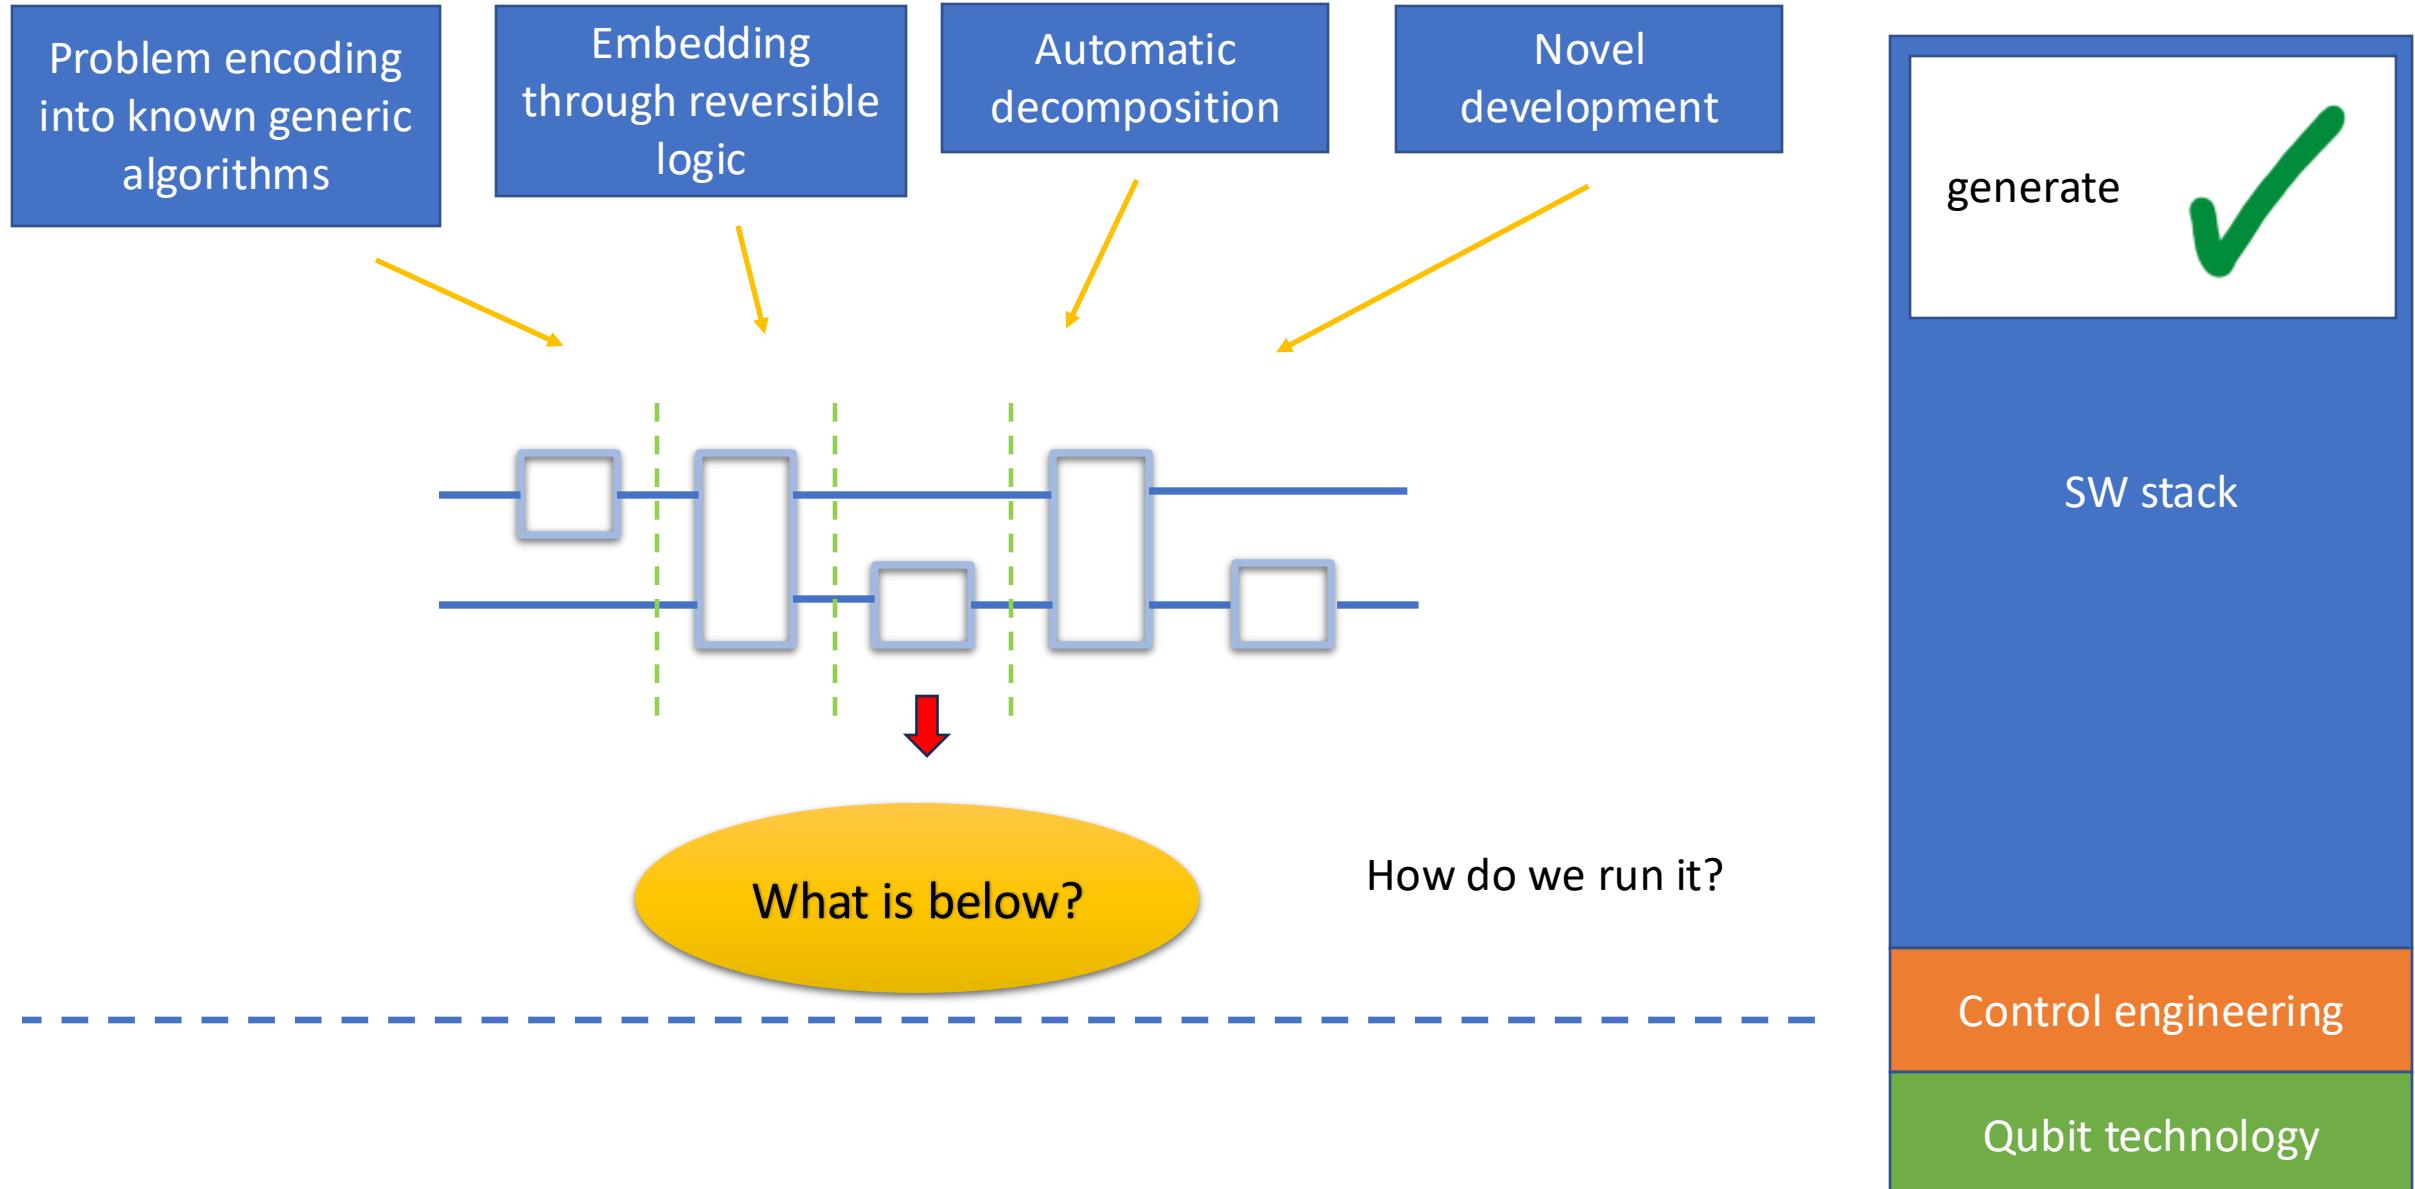
Matrix form

$$F_N = \frac{1}{\sqrt{N}} \begin{bmatrix} 1 & 1 & 1 & 1 & \cdots & 1 \\ 1 & \omega & \omega^2 & \omega^3 & \cdots & \omega^{N-1} \\ 1 & \omega^2 & \omega^4 & \omega^6 & \cdots & \omega^{2(N-1)} \\ 1 & \omega^3 & \omega^6 & \omega^9 & \cdots & \omega^{3(N-1)} \\ \vdots & \vdots & \vdots & \vdots & & \vdots \\ 1 & \omega^{N-1} & \omega^{2(N-1)} & \omega^{3(N-1)} & \cdots & \omega^{(N-1)(N-1)} \end{bmatrix}$$

4. Novel design

- ❖ Not an easy task
- ❖ Much of our reasoning is still tied to circuits and complex Hilbert spaces
- ❖ We are “chasing vectors around” in an analogy to “chasing bits around”
- ❖ Very active fields in quantum algorithm theory are:
 - Quantum error correction codes
 - Quantum complexity classes
 - MIP* = RE, Certifiable randomness,
Classically verifiable quantum advantage
 - Finding new classical algorithms by “dequantization”

Gate-based quantum computing model



A number of circuit optimizations

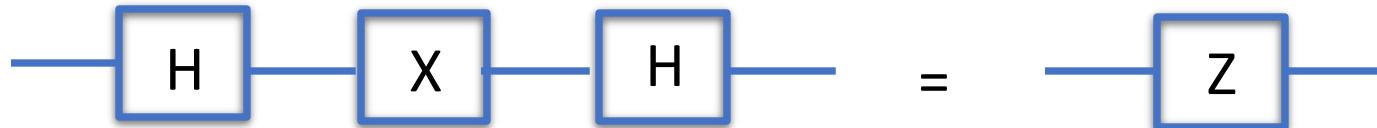
- ❖ **Circuit compression** – minimize the number of gates used (focus on coupling gates in particular)
- ❖ **Unroll/decompose** to the native gate set supported by the quantum HW
- ❖ **Optimal routing** – map the logical circuit to the physical chip while respecting its connectivity map. Insert SWAP gates where needed.
- ❖ (Insert **error mitigation** gates).

These optimizations techniques are partwise orthogonal, quantum HW dependent, and may be applied iteratively/recursively in order to achieve the best results.

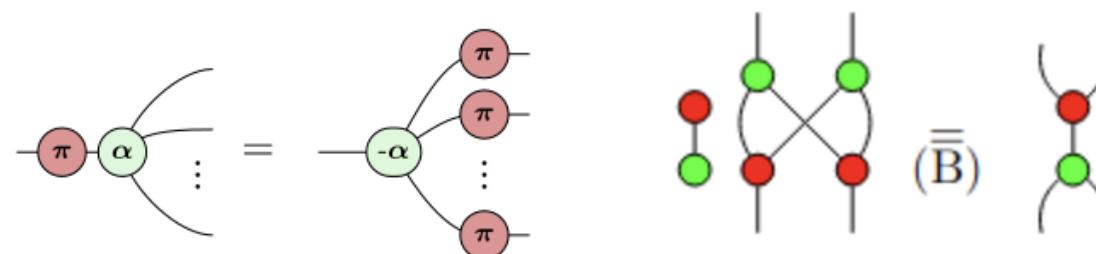
Circuit compression

- ❖ The most common technique is to exploit **logical circuit identities**

Eg:



- ❖ One of the newer approaches is called **ZX-calculus**.
 - It relaxes the unitarity condition: operates in a less restrictive linear regime instead
 - But, it's not always possible to revert back to a unitary circuit



Unrolling/decomposition

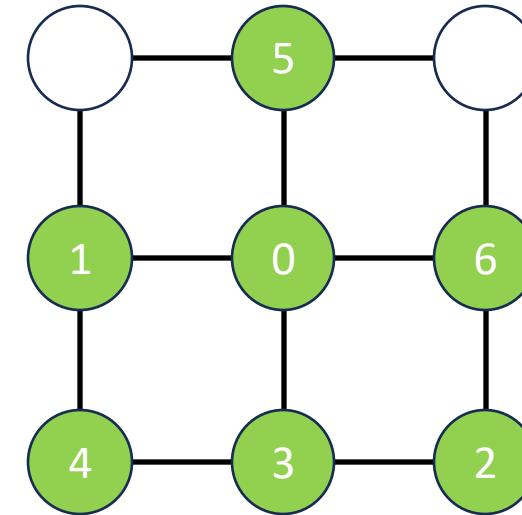
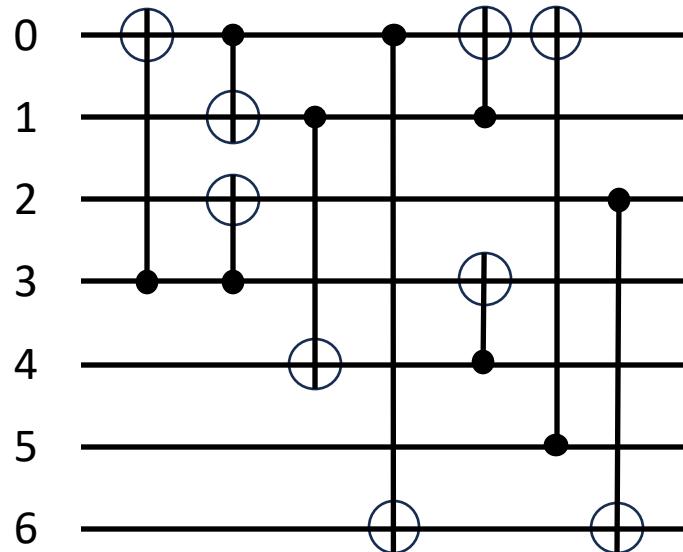
- ❖ There are many universal gate sets for quantum computing.
- ❖ For superconducting qubits, common entangling gates are: **CX**, **CZ**, or **iSWAP** accompanied with **Rx(..)** and **Rz(..)** single qubit rotations. We call it a **native** gate set.
- ❖ SW stack typically contains a **library** of definitions of other **common gates** in terms of the native universal gate set. Thus, for example, the Hadamard gate H can be ‘*unrolled*’ in terms of Rx(..) and Rz(..) as:



- ❖ Uncommon gates needs to be (brute-force) decomposed (eg. by SVD).

Optimal routing

- ❖ A superconducting quantum chip typically supports only interactions between nearest-neighbour qubits. We talk about a connectivity map.
- ❖ More distant interactions are achieved via inserting (multiple) SWAP gates. We want to minimize the number of burdensome SWAPs.



This problem is quite similar to a CPU register allocation.

L3 cache

L2 cache

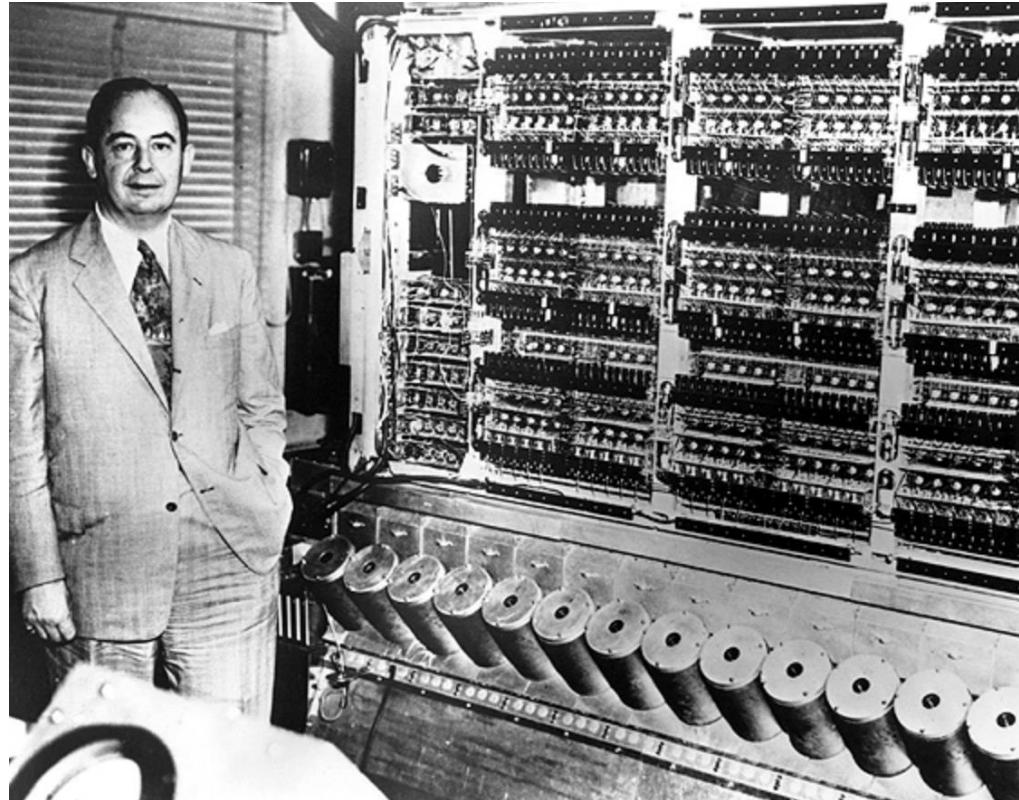
Intermezzo: Connectivity map & memory

Intel Core i9-13900K with 8 P-cores and 16 E-cores

- Caches (all levels, 68MB) $\approx 4.1 \times 10^9$ transistors.
- All integer ALUs (whole chip) $\approx 1 \times 10^6$ transistors.
- FP/vector units (if included) push compute logic into $\sim 10^6\text{--}10^7$ range.



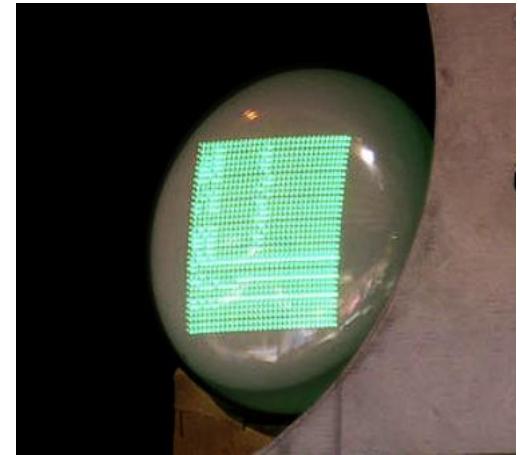
World's first random-access memory at IAS



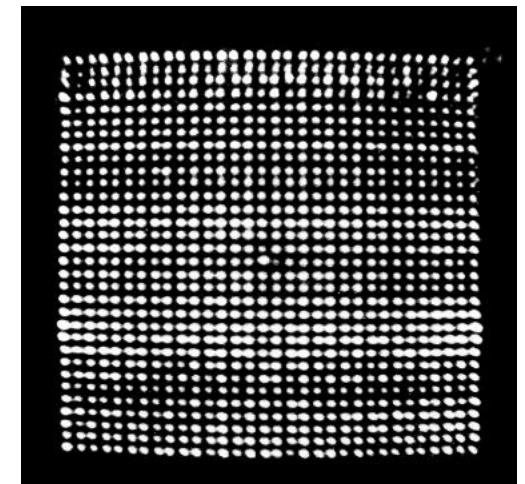
John von Neumann



James Pomerene



32x32 CRT

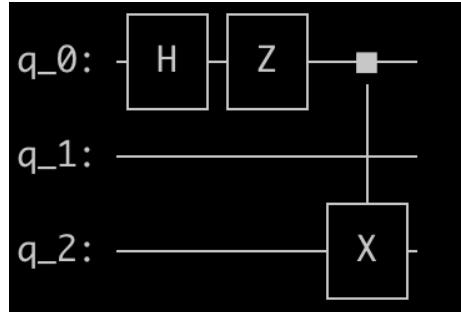


Diagnostic photograph
from maintenance logs

Example: Qiskit's built-in circuit optimizations

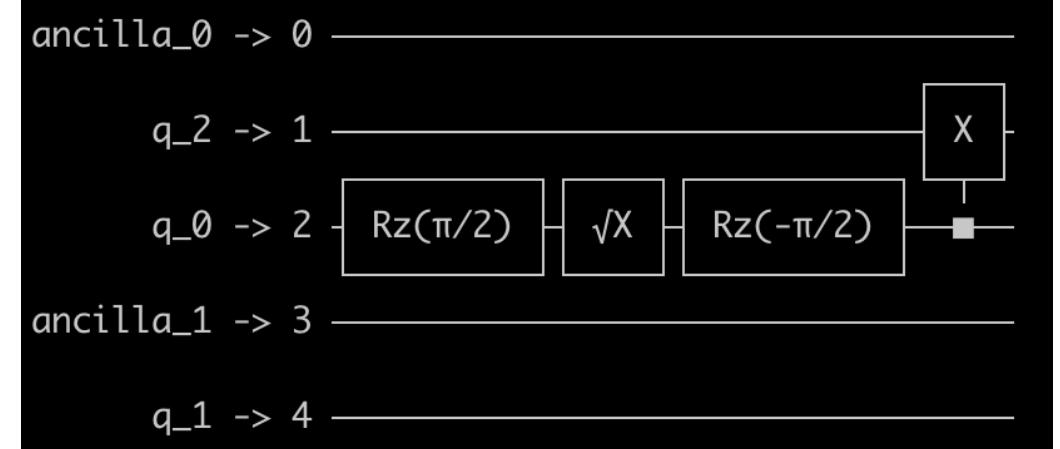
Original circuit

```
circ = QuantumCircuit(3)
circ.h(0)
circ.z(0)
circ.cx(0,2)
```



Transpiled circuit

```
trc = transpile(circ, backend)
```

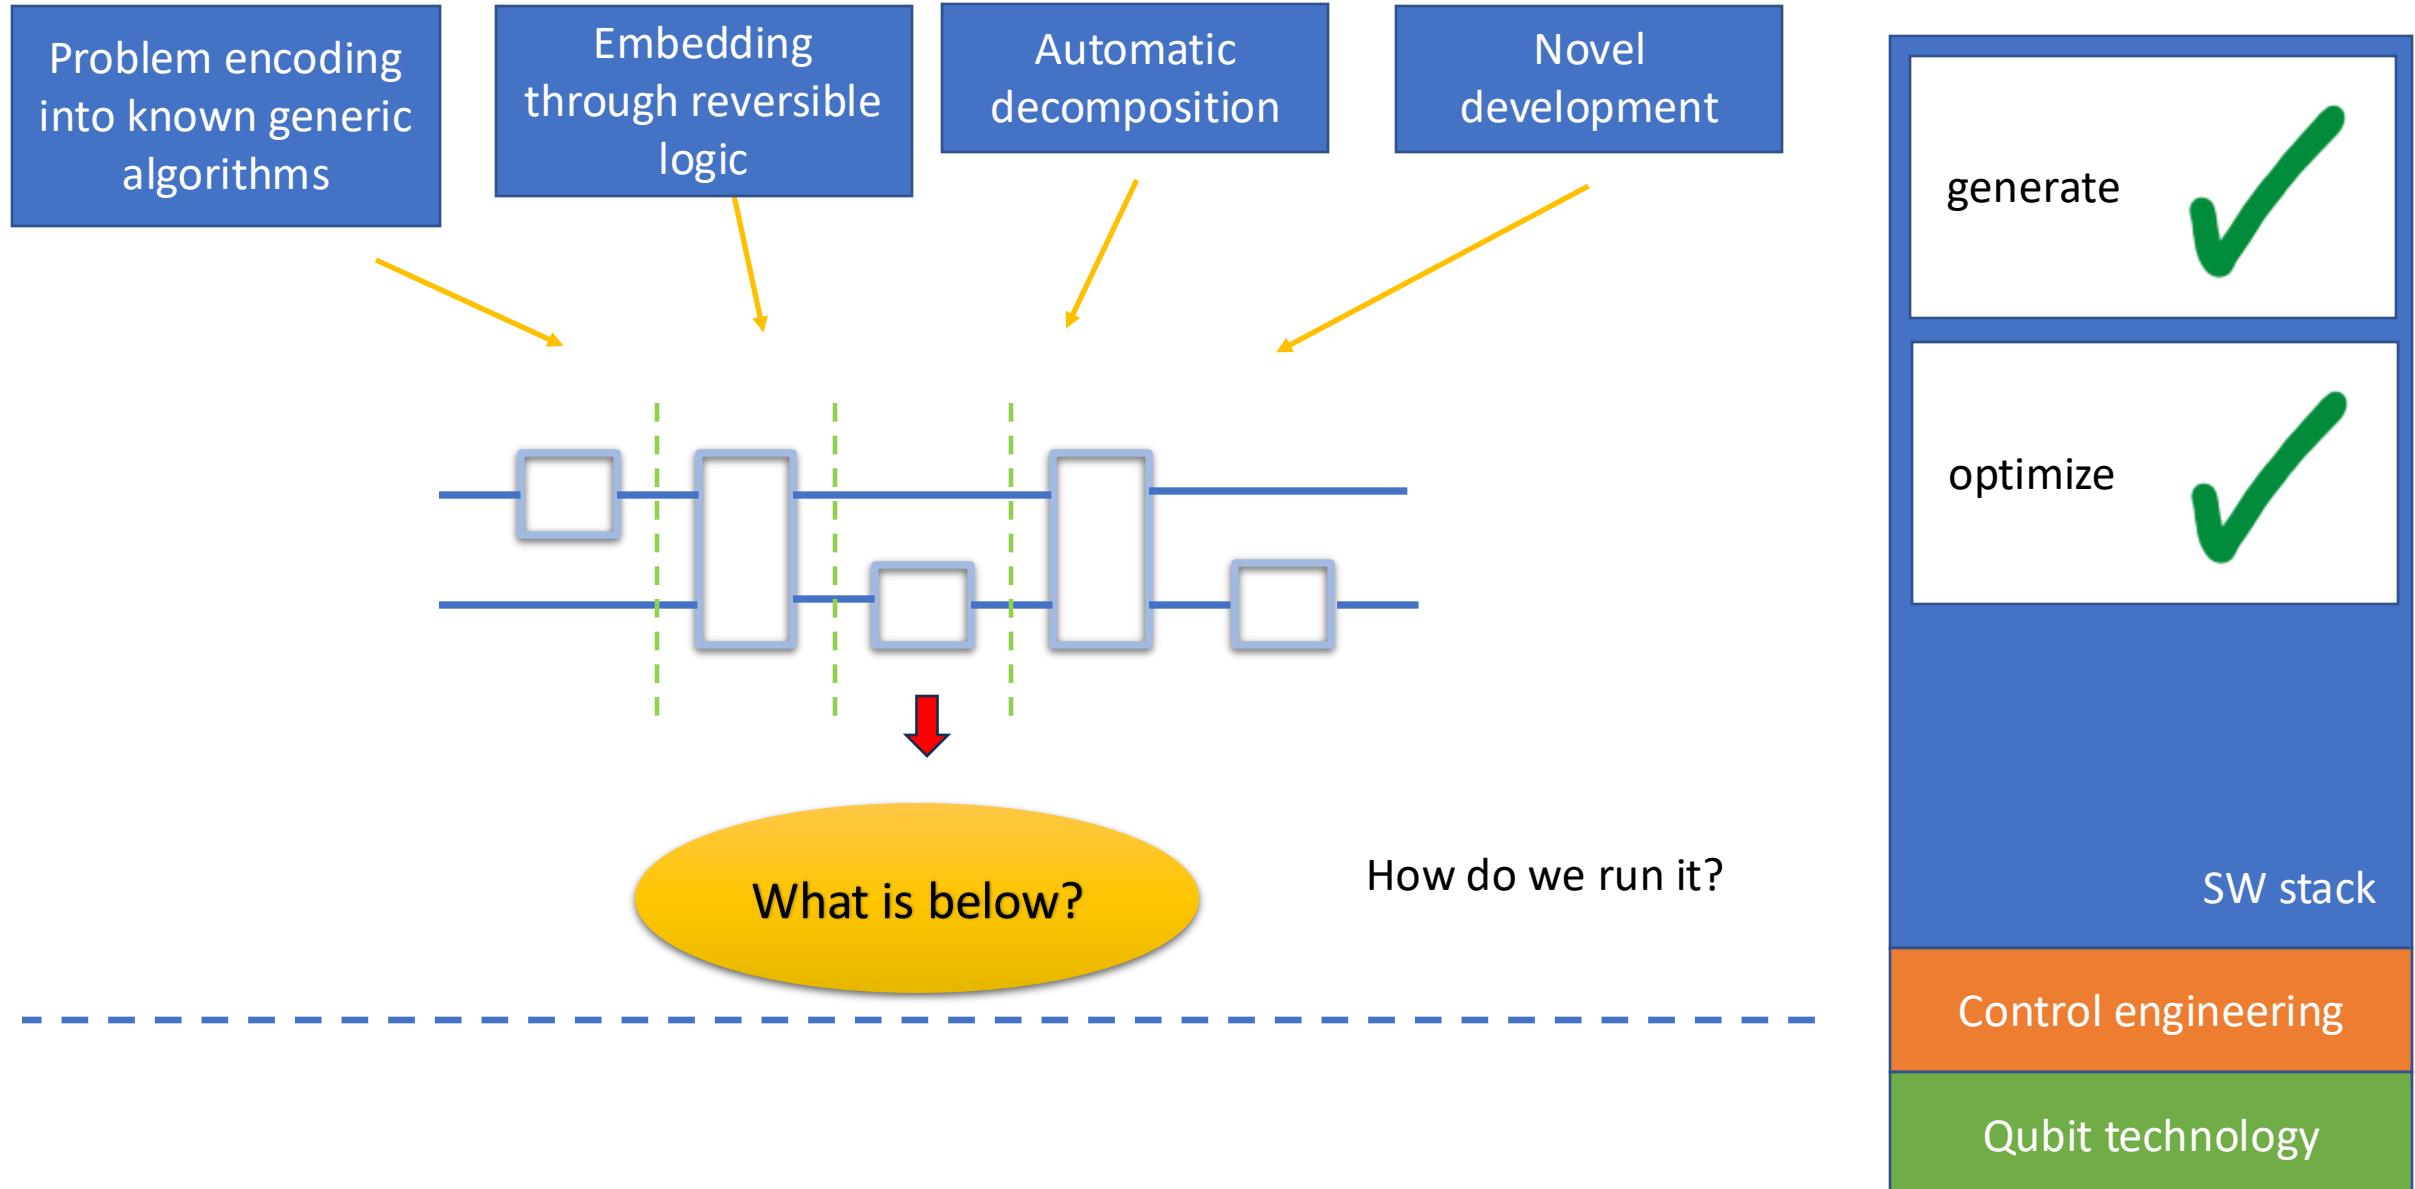


Manila's coupling map



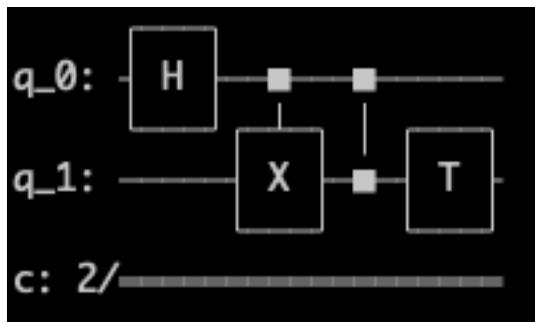
Unrolling, compression
and routing has been applied.

Gate-based quantum computing model



Quantum circuit execution

- ❖ The generated & optimized circuit needs to be converted from an internal high-level representation (say a Python object) to a flattened textual or binary representation suitable for network transfer and execution.

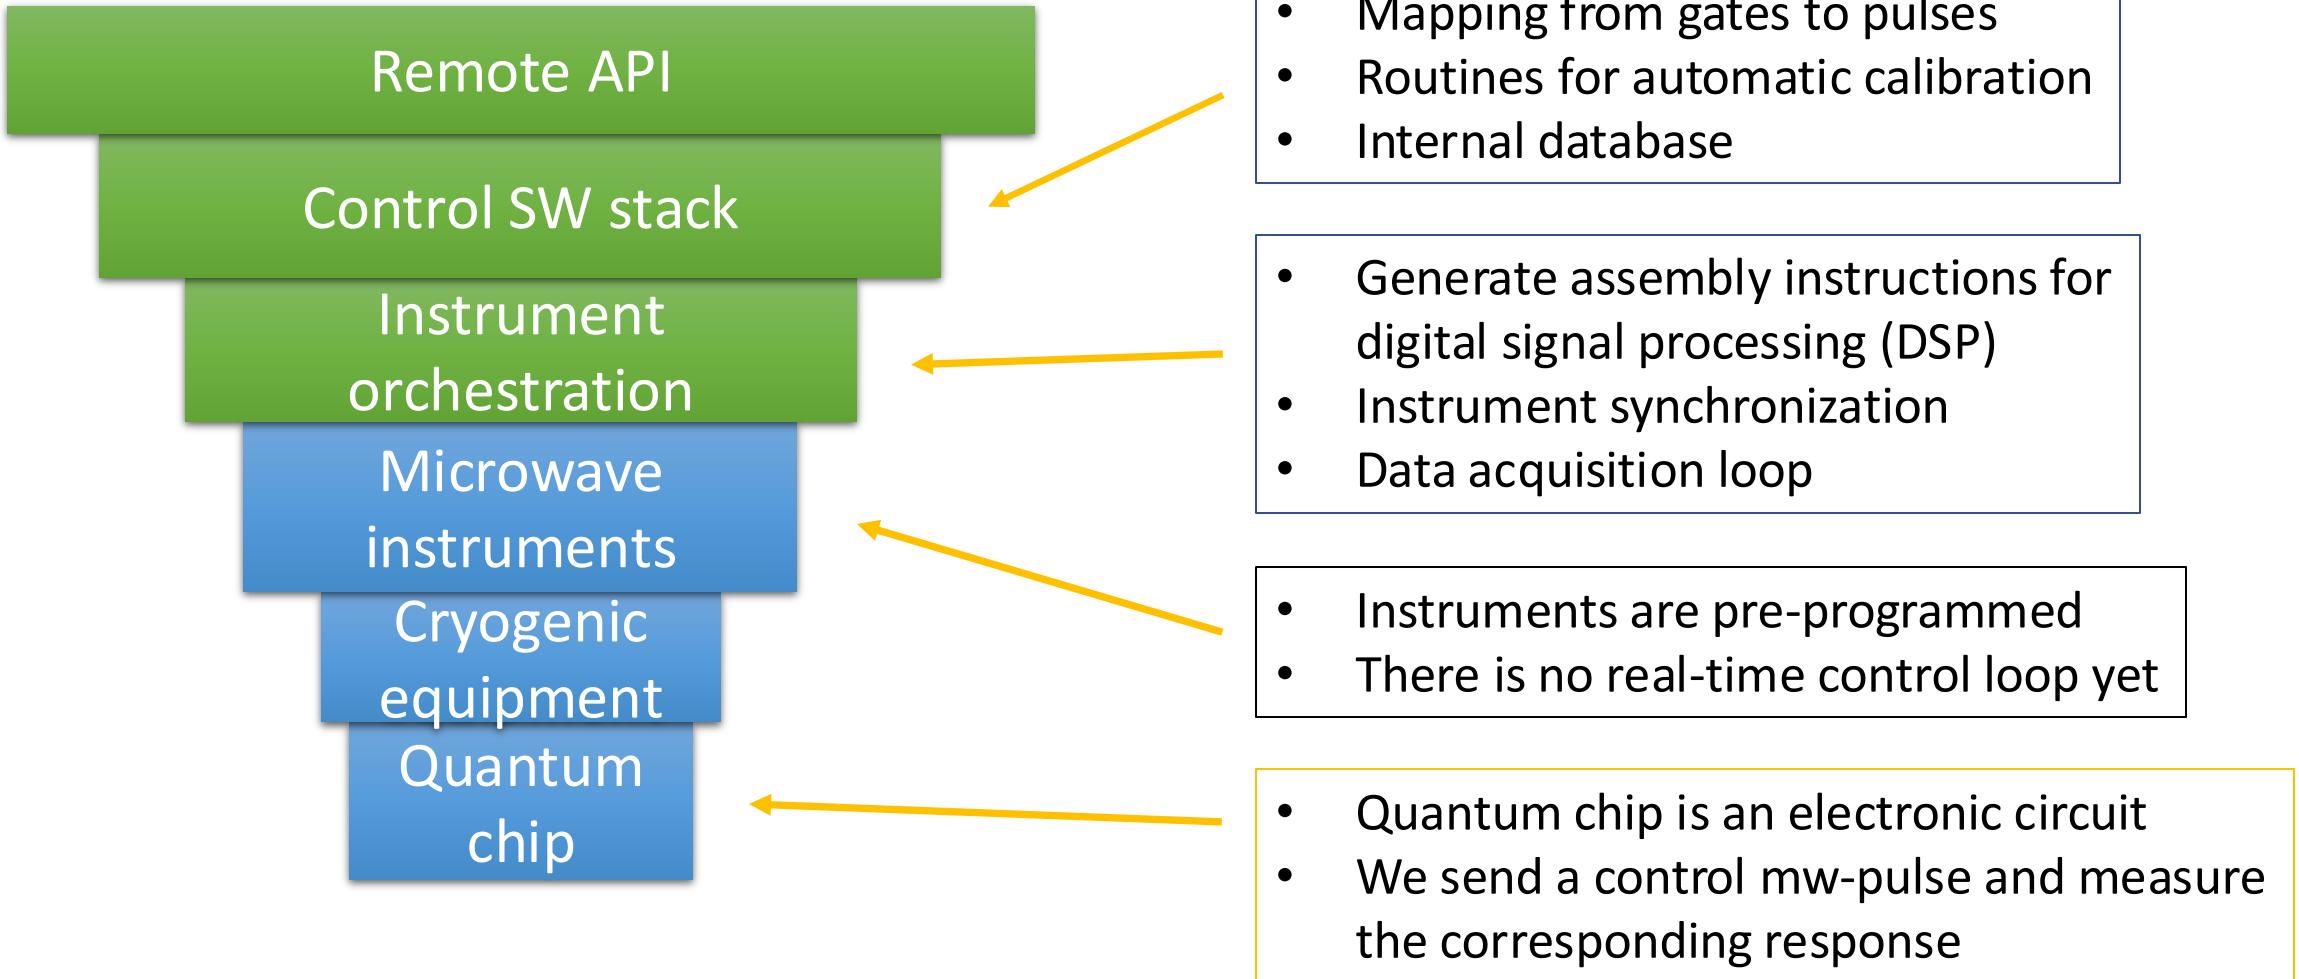


assemble
→

```
OPENQASM 2.0;
include "qelib1.inc";
qreg q[2];
creg c[2];
h q[0];
cx q[0],q[1];
cz q[1],q[0];
t q[1];
```

- ❖ *OpenQASM v2* from IBM has emerged as a practical standard due to its simplicity and permissive licensing.
- ❖ *OpenQASM v2* is also often used as inter-operability language between different circuit toolkits.

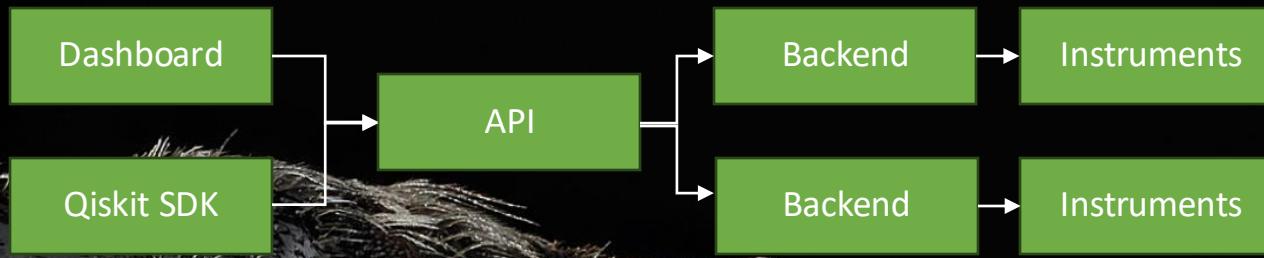
Execution target: NISQ device



Tergite: software stack



- ❖ Cloud service for the quantum computer
- ❖ Automated chip bring-up
- ❖ Integrated with Puhuri
- ❖ Open-source code available on GitHub
 - ❖ Apache 2.0
 - ❖ <https://tergite.github.io/>

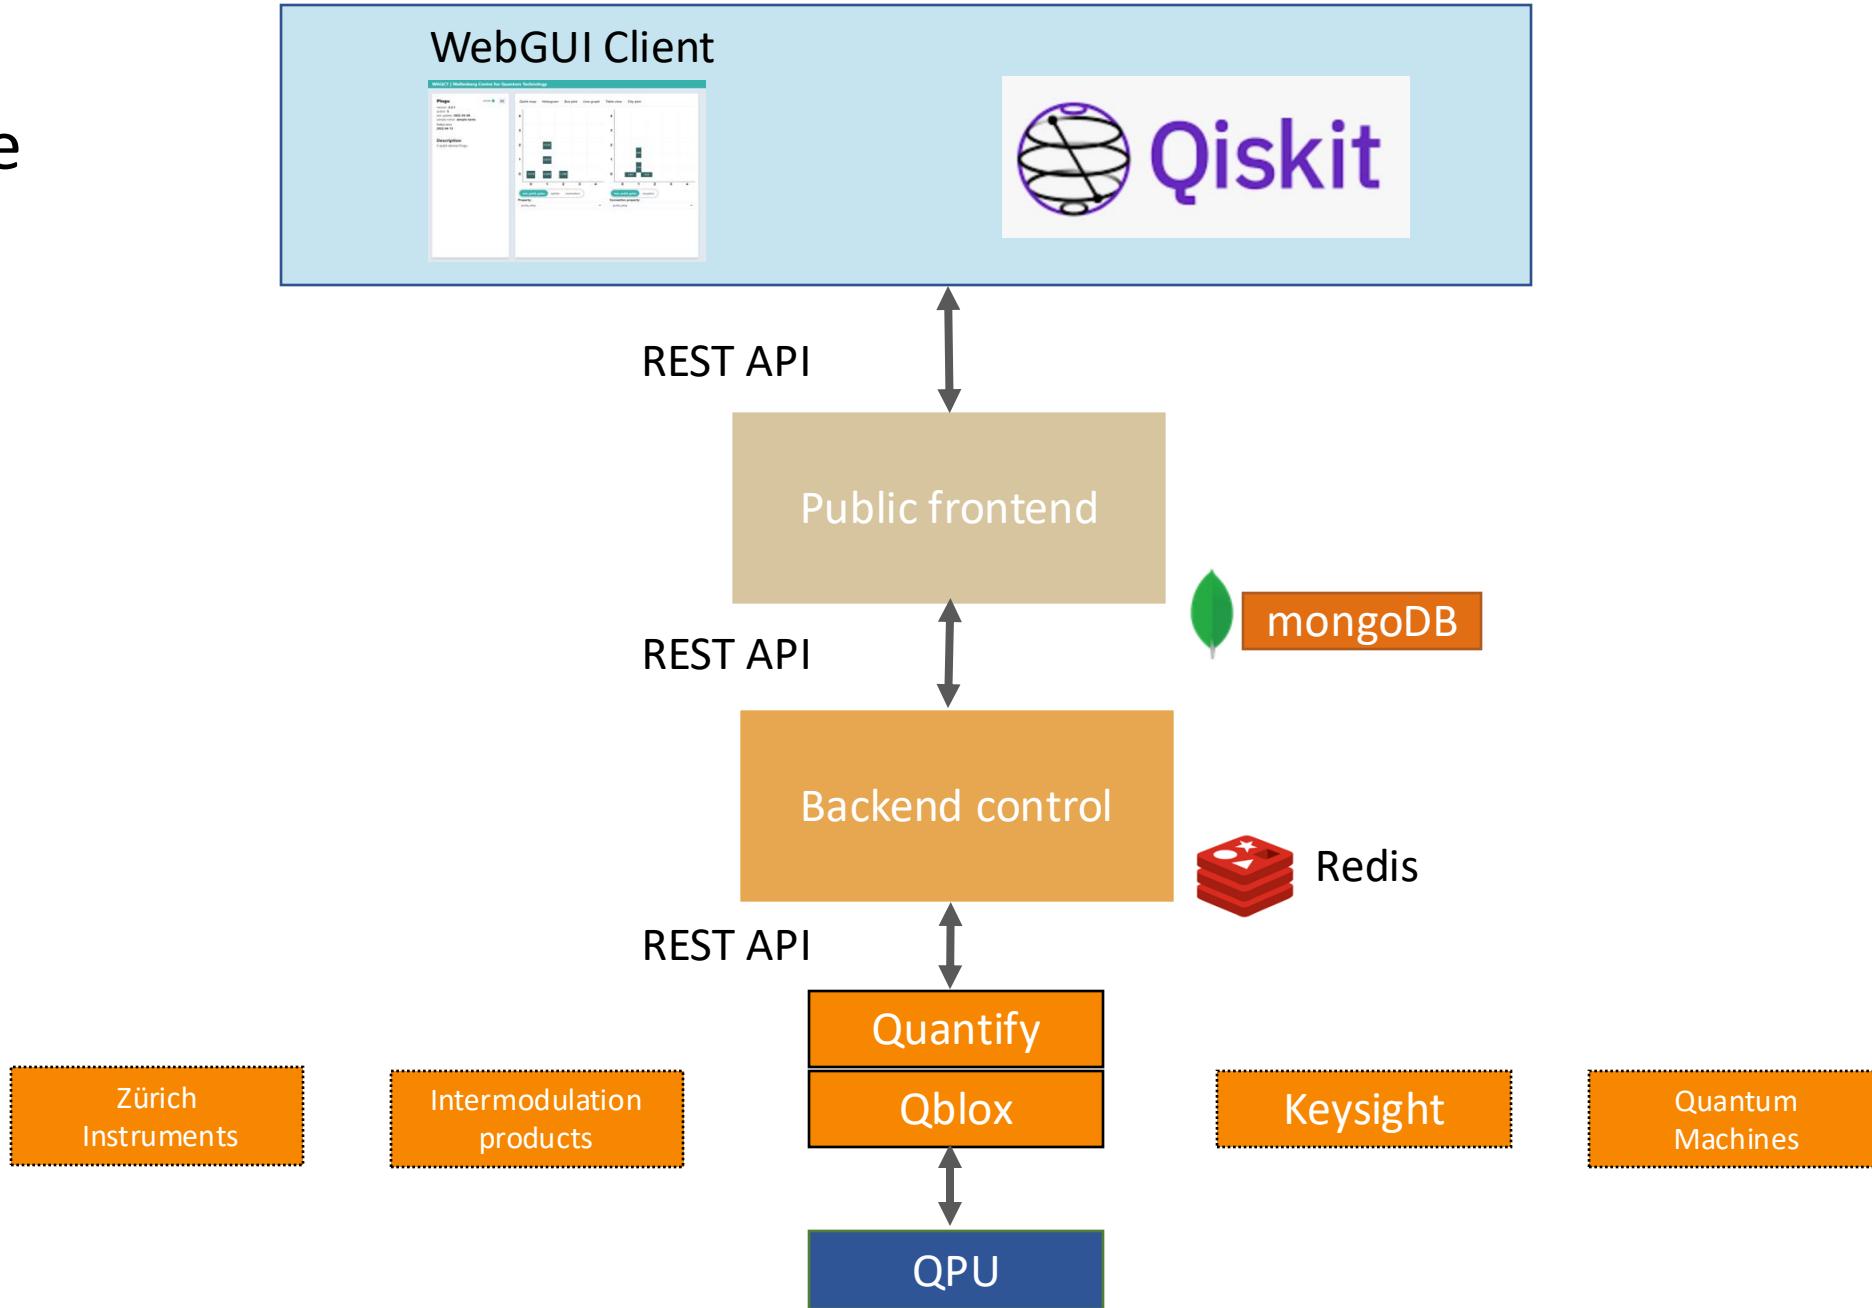


The screenshot shows the GitHub profile for the 'tergite' organization. It features a large profile picture of the Tergite logo. The 'Overview' tab is selected, showing 8 followers and a location in Sweden. Three projects are pinned to the top of the repository list:

- tergite-autocalibration** (Public): A commandline application to calibrate the WACQT quantum computers automatically. It includes a Jupyter Notebook and has 3 stars and 4 forks.
- tergite-frontend** (Public): The Frontend in the [Tergite software stack] (<https://tergite.github.io/>) of the WACQT quantum computer. It is implemented in Python and has 1 star and 1 fork.
- tergite-backend** (Public): The Backend in the Tergite software stack of the WACQT quantum computer. It is implemented in Python and has 3 stars and 3 forks.

On the right side of the GitHub interface, there is a preview of the Tergite dashboard, which displays various performance metrics and graphs related to the quantum computer's calibration process.

Tergite





Qibo · v0.2.21

Search

INTRODUCTION

[Getting started](#)

[Code examples](#)

MAIN DOCUMENTATION

[API reference](#)

[Developer guides](#)

APPENDIX

[Publications](#)

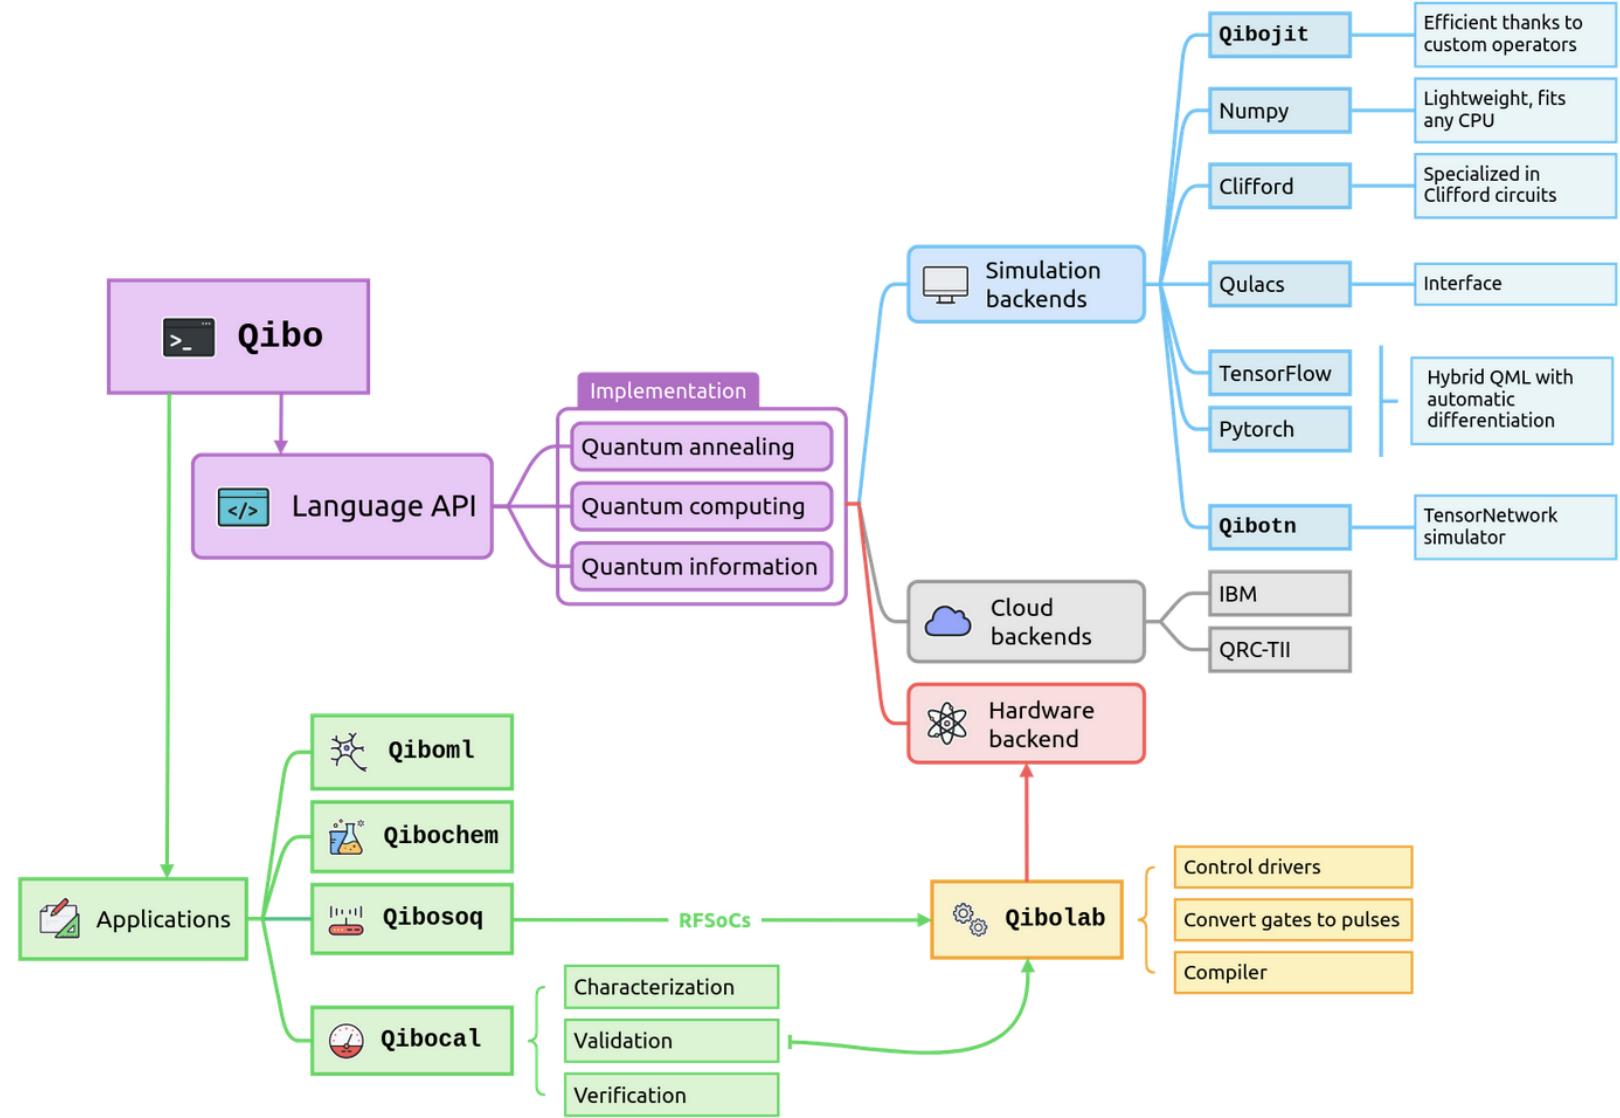
DOCUMENTATION LINKS

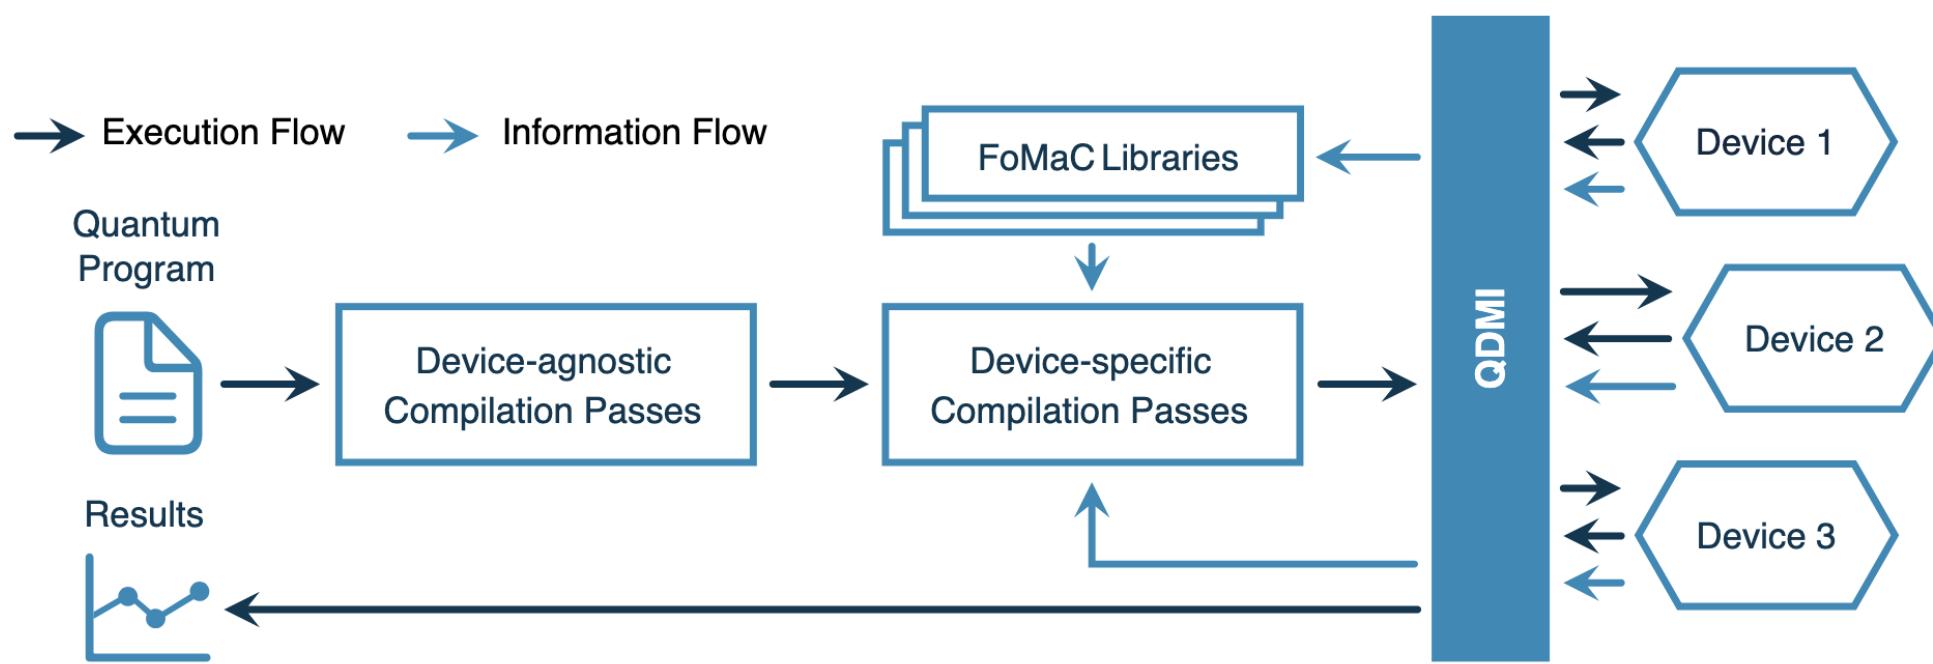
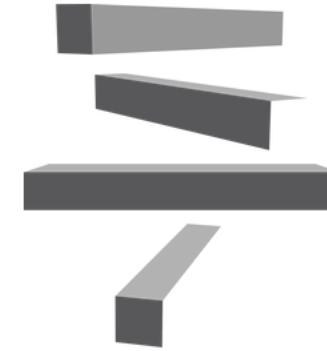
[Qibo docs](#)

[Qibolab docs](#)

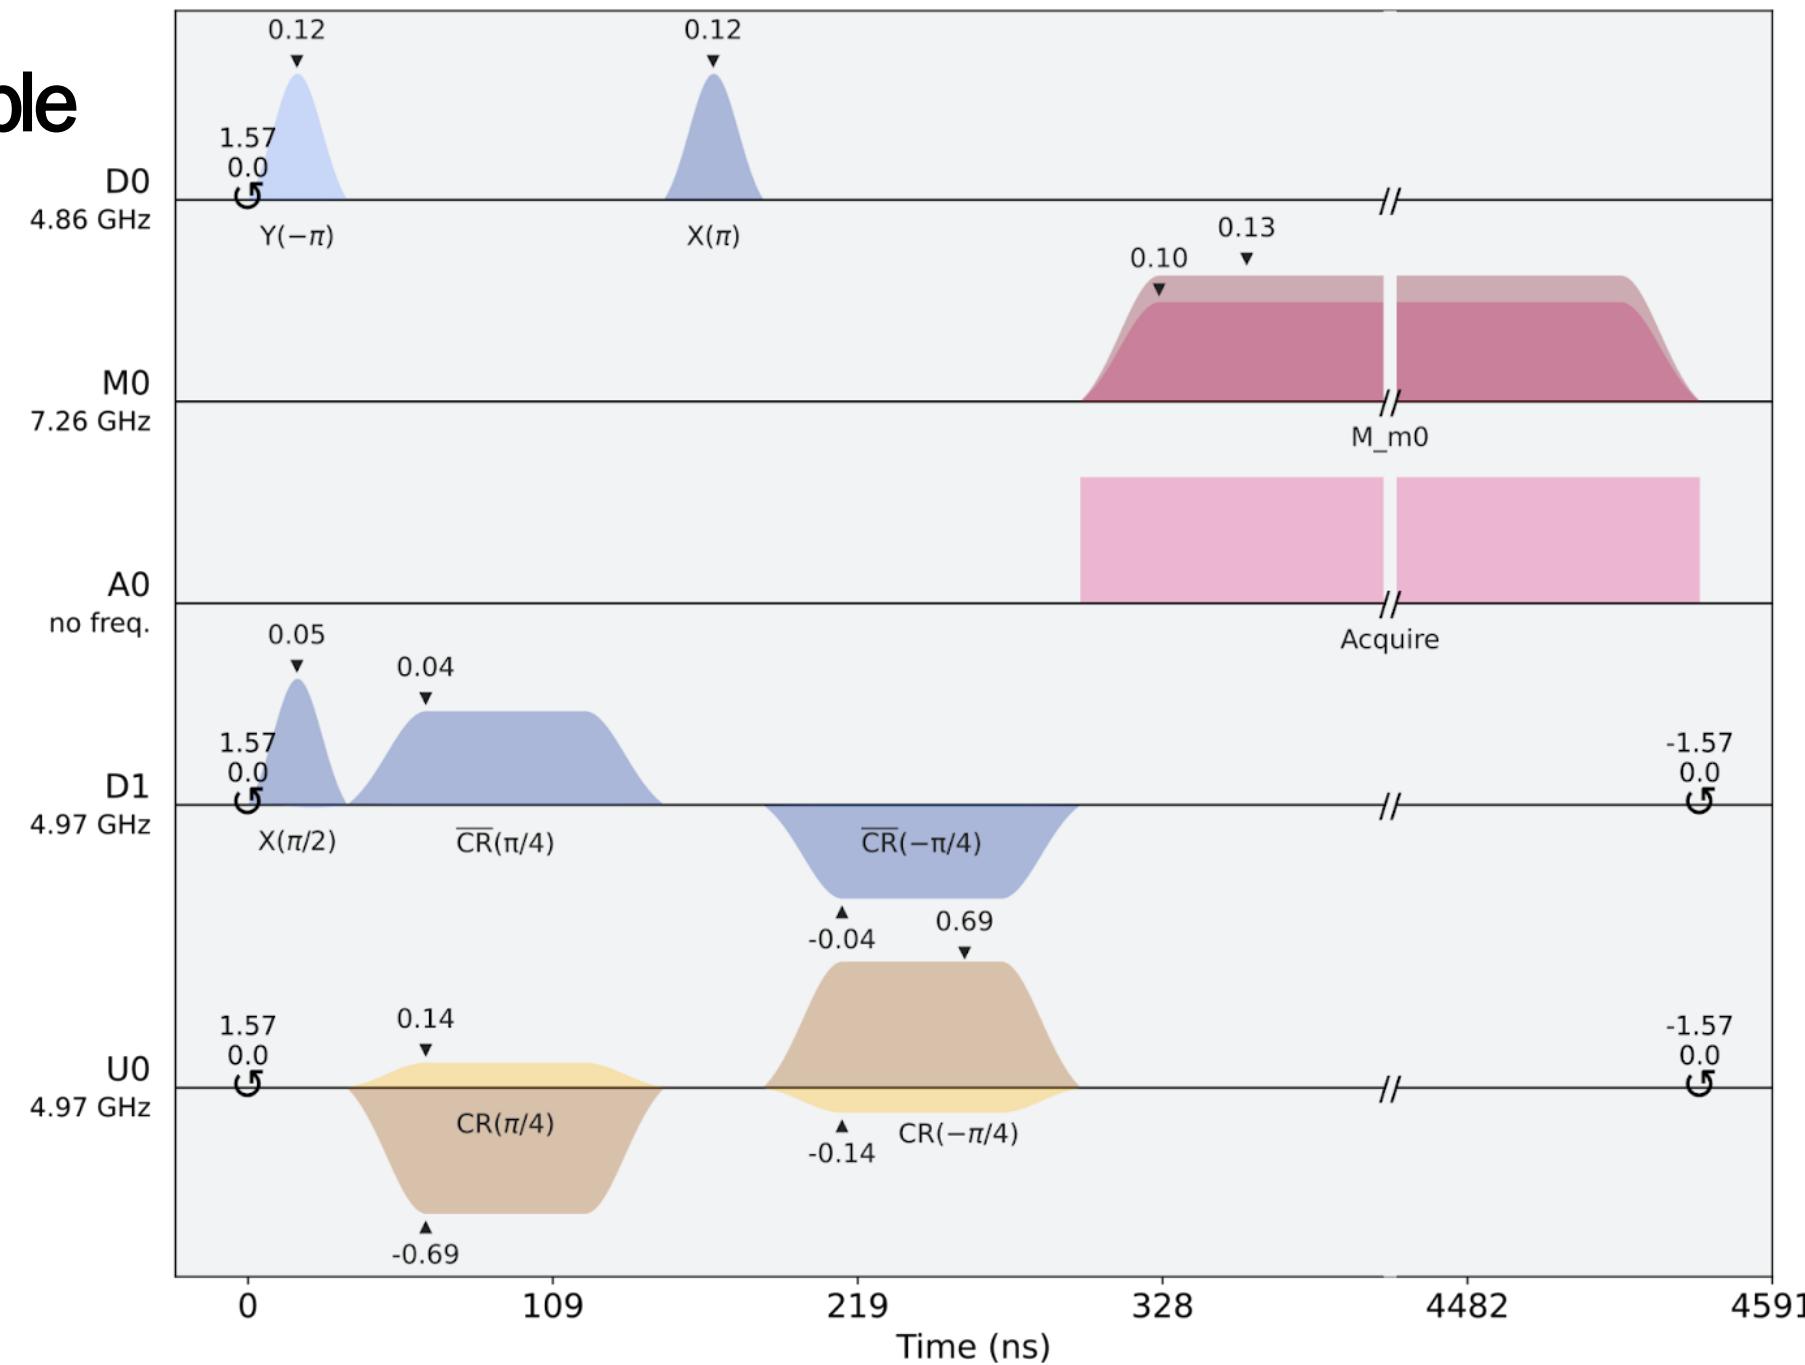
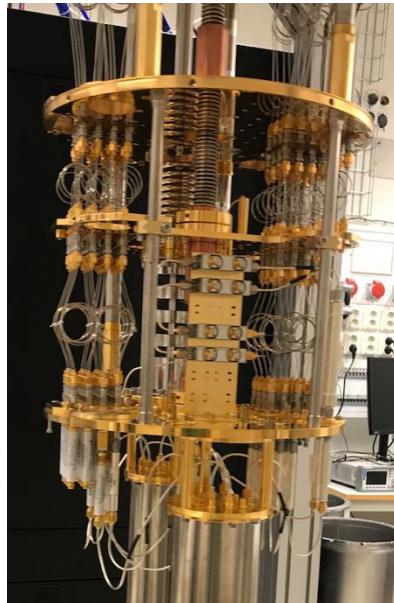
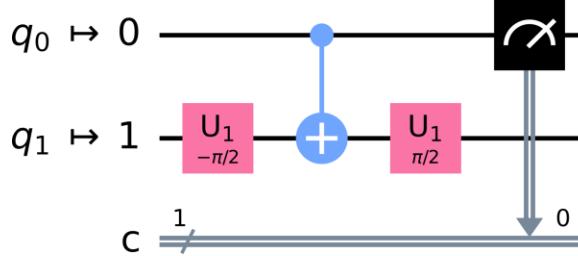
Components

The main components of Qibo are presented in [Getting started](#)





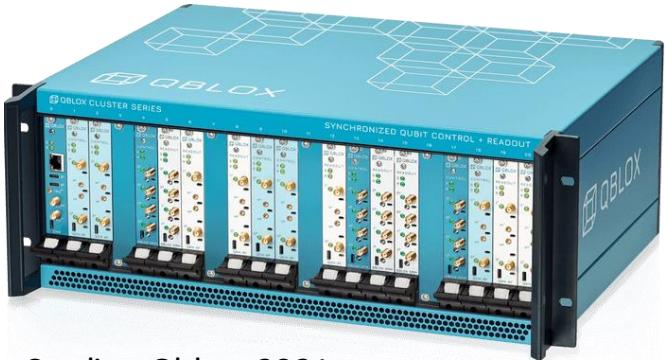
Pulse schedule example



Example: Qblox instruments assembly

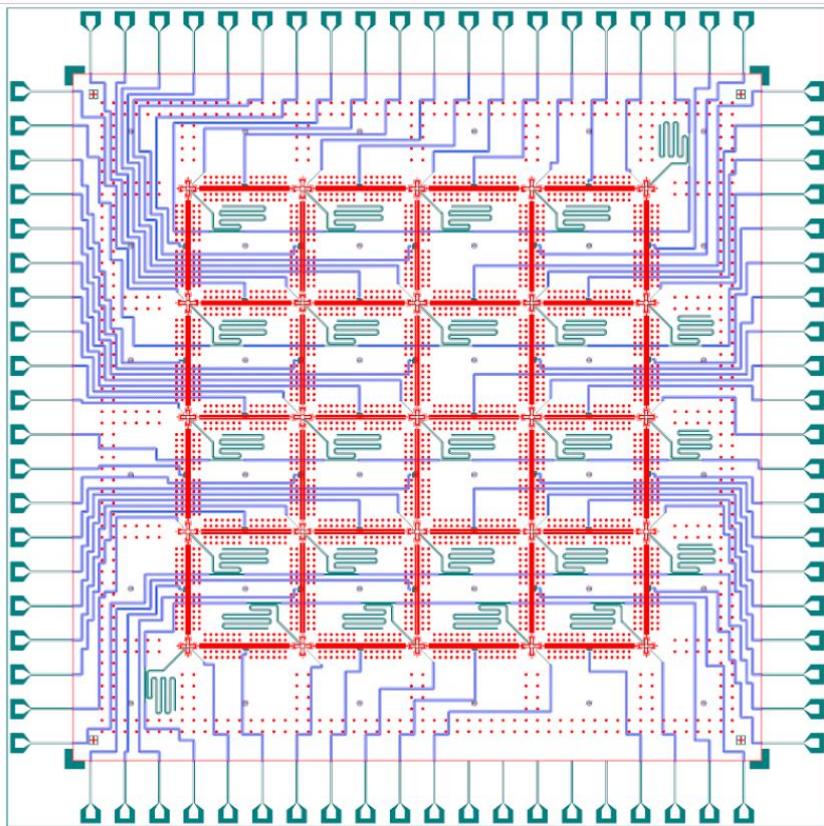
Q1ASM program:

```
0:          wait_sync      4
1:          upd_param     4
2:          set_mrk       15      # set markers to 15
3:          wait           4      # Latency correction of 0 ns.
4:          move          2000,R0 # iterator for loop with label start
5:          start:
6:          reset_ph
7:          upd_param     4
8:          wait           65532   # auto generated wait (300000 ns)
9:          wait           65532   # auto generated wait (300000 ns)
10:         wait          65532   # auto generated wait (300000 ns)
11:         wait          65532   # auto generated wait (300000 ns)
12:         wait          37872   # auto generated wait (300000 ns)
13:         set_awg_gain  851,0    # setting gain for gaussian-d1-0
14:         play          0,1,4    # play gaussian-d1-0 (100 ns)
15:         wait           96      # auto generated wait (96 ns)
16:         wait           4       # auto generated wait (4 ns)
17:         set_awg_gain  851,0    # setting gain for gaussian-d1-104
18:         play          0,1,4    # play gaussian-d1-104 (100 ns)
19:         wait           3596    # auto generated wait (3596 ns)
20:         loop          R0,@start
21:         set_mrk      0        # set markers to 0
22:         upd_param     4
23:         stop
```



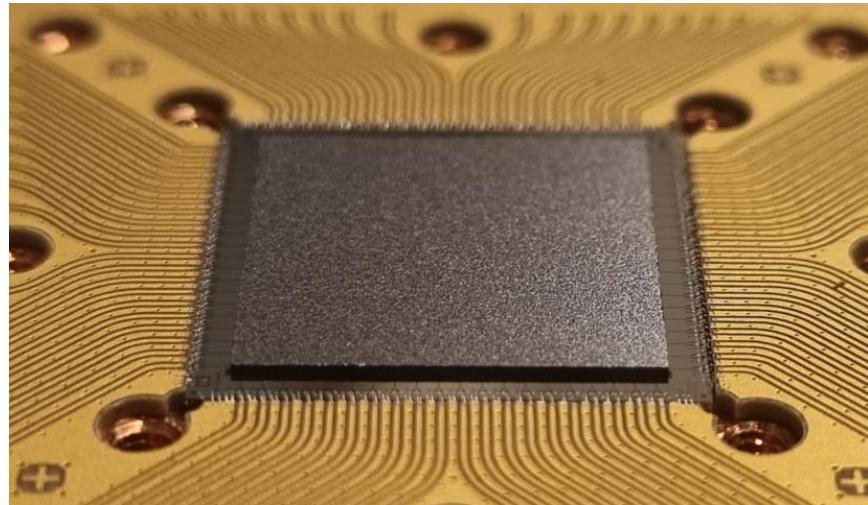
Credits: Qblox, 2021

QPU chip

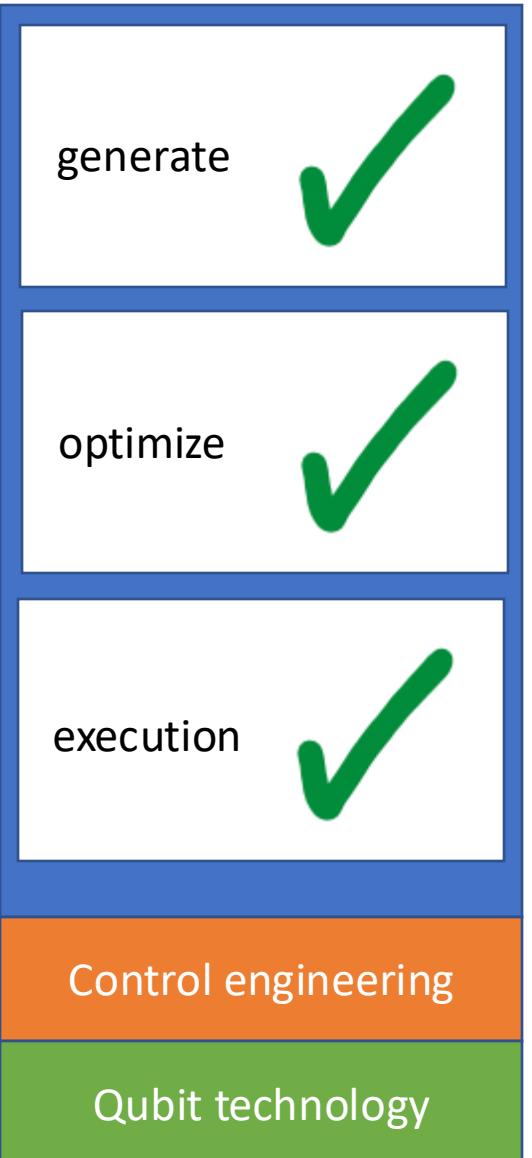


A layout of a 25 qubit processor developed at Chalmers.

WACQT

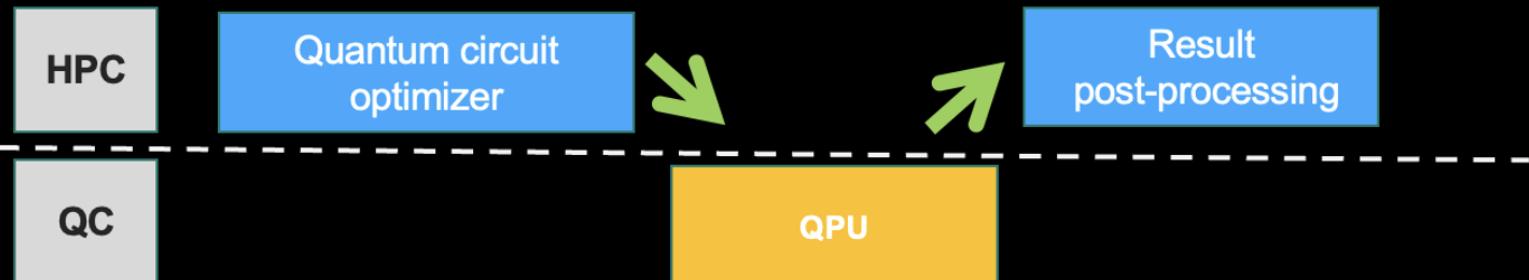


■ Substrate	■ Bump	■ Qubit & Coupler
■ Ground	■ UBM	■ Control & Readout

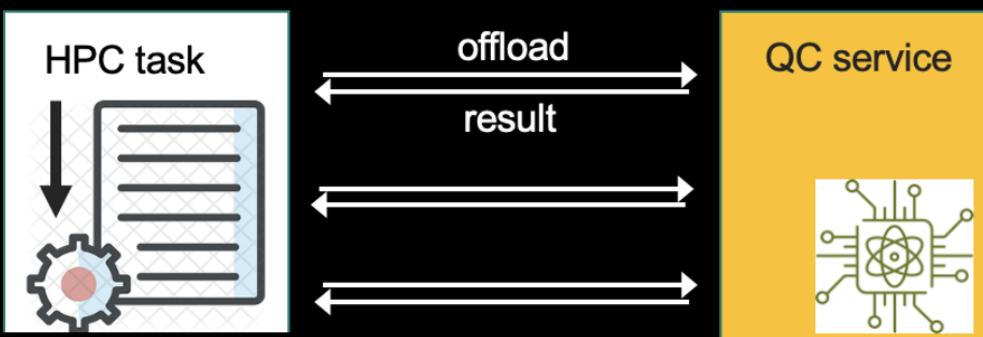


HPC-QC use-cases

1) Using HPC power in pre-processing and post-processing of quantum jobs



2) Using QC as a special purpose accelerator of sort (sampling from a distribution)



Monte Carlo simulation
estimating the number PI

